

Today - how could talloc have been better?

- Talloc does garbage collection!

A gc is supposed to free up memory that's not being used anymore. (C doesn't have one)

↙ take time and memory
unavoidable,
but maybe worth it

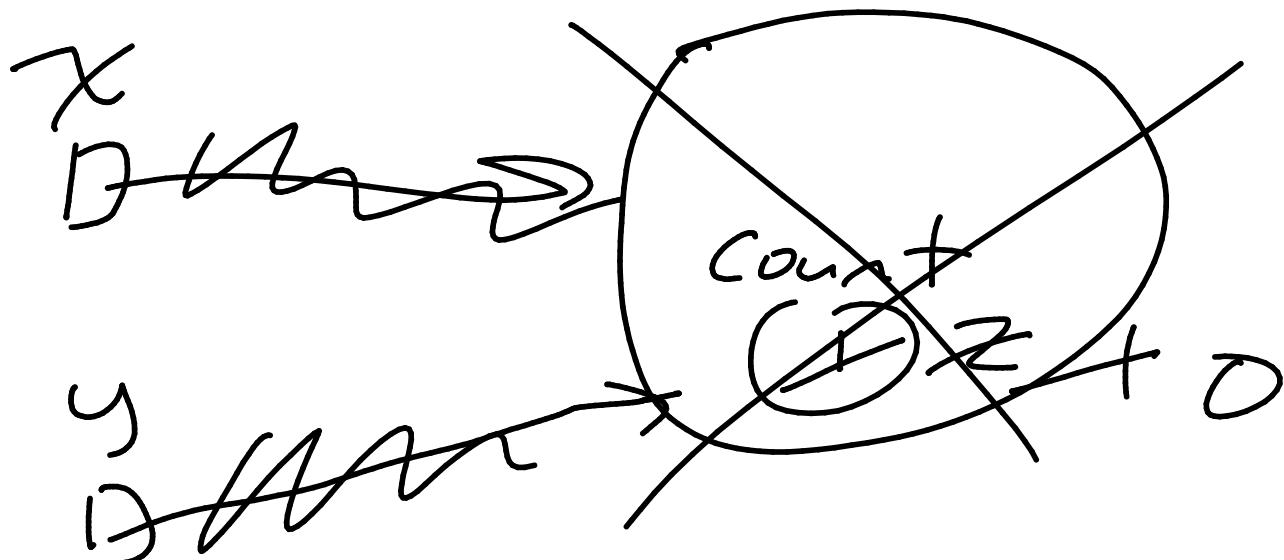
How is it done better?

① straightforward, good in some ways,
and somewhat common, but also
doesn't quite work

Reference Counting

Every object [SchemeVal] has a count
of the number of pointers referring
to it.

When you assign a new one, increment
- - move one away, decrement
when counter is 0, you free it



$y = x$

$y = \text{null}$

$x = \text{null}$

What's good?

Easy

Incremental - frees
memory as soon as it
becomes free

Relatively fast

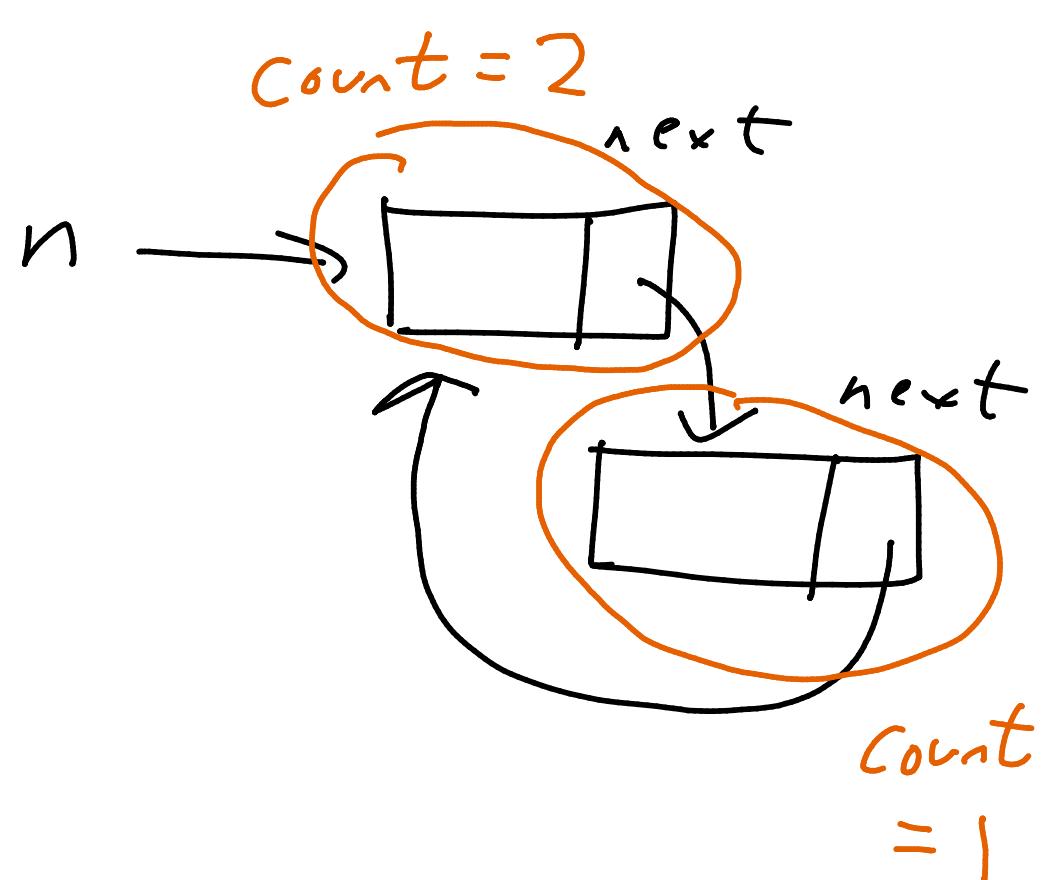
What's the problem?

- circular references

$n = \text{Node}()$

$n.\text{next} = \text{Node}()$

$n.\text{next}.\text{next} = n$

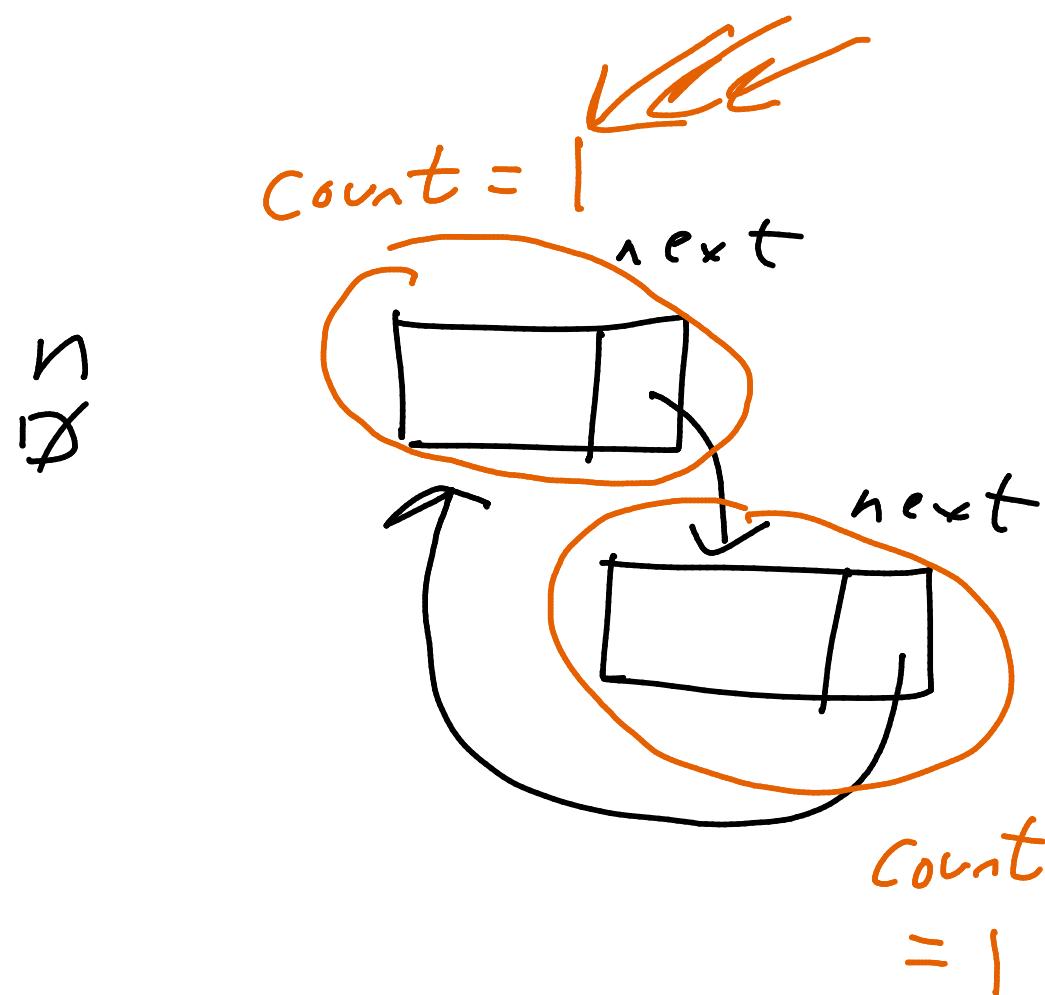


$n = \text{null}$

Count == 1,

so nothing
gets freed

- but it's all
garbage!



How do you handle this?

???

① You don't, ugh

② You fit some kind of ugly
cycle detection algorithm on
top of this

C Python (python.org)
does this. Ref counting +
occasional cycle detecting

② (actually does work)

Mark-and-sweep

- occurs "every so often"

Follow all pointers that you have
access to (in your function call
stack, e.g. by all variables
that are live)

- mark those objects as alive
recursively, mark all the objects
that they point to, and so on.

When done marking, free all
the unmarked objects.

Accuracy?

Time?

Memory needed?

Other?

Where did I oversimplify that would require effort / time / memory / etc?

* How do you find the dead objects?

When are you doing this?

Tradeoff of time spent vs free memory

Typically

- time intervals
- memory gets "close" to full

How do you find all the pointers?

- You need some kind of clear object defns and known pointer types within?

How do you find all current variables?

In our Scheme implementation,
look in current frame, its
parent, grandparent, etc.

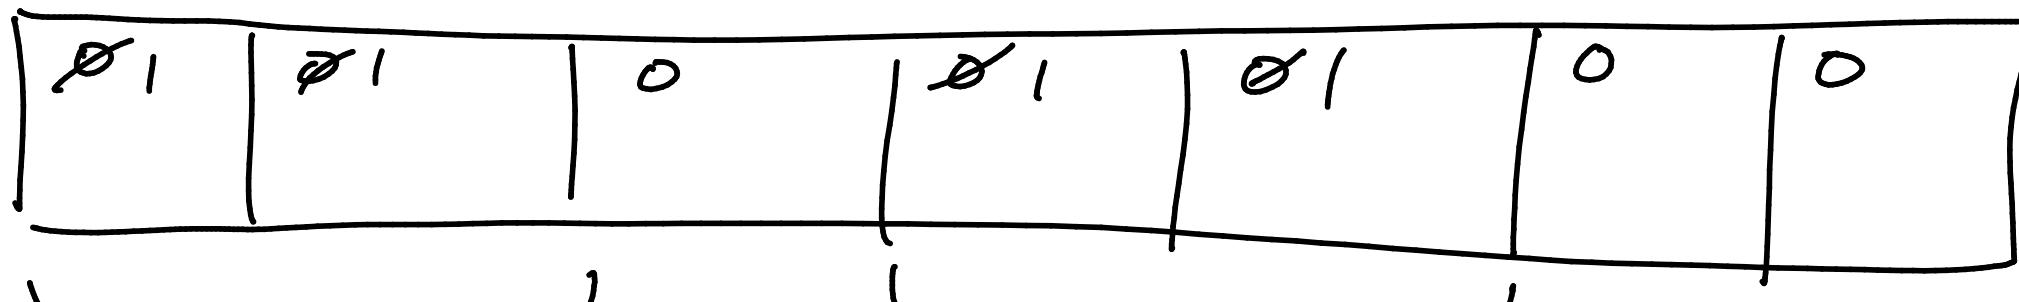
[follow everything recursively]

This does correctly handle circular
references.

* How find dead objects?

The heap is broken up into
small fixed-size units, which
are combined to make bigger units
if you need them for your objects.

- there's a space in each unit
indicating if it was allocated
when the object was constructed



get me a int

set me a double

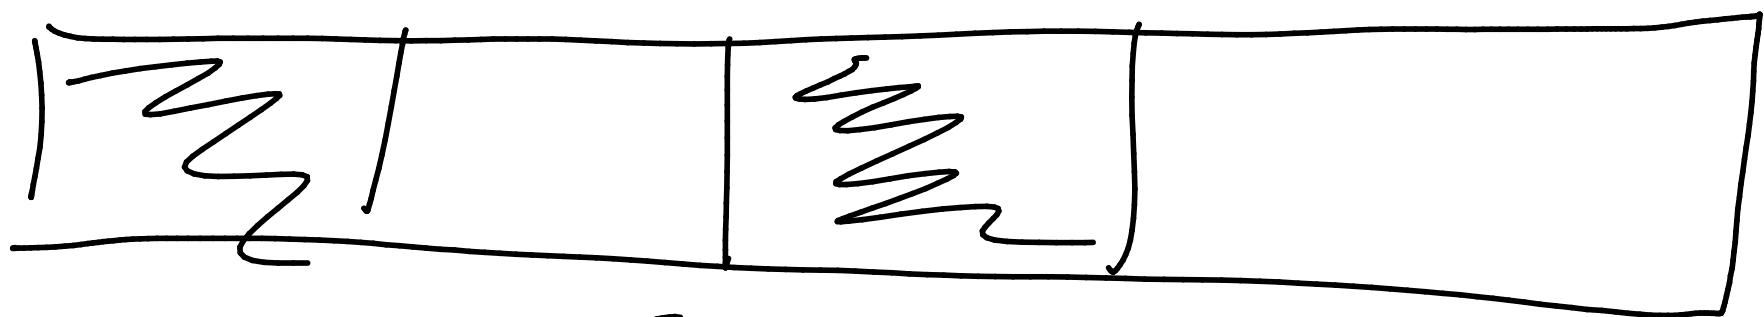
this is not marking from gc, this
is just an indication if it had
ever been allocated

To find dead objects, loop over
entire heap. If it had been
allocated at some point and
is not marked, free it.

One more big issue is that this
approach leaves memory fragmented



↓ free



↑

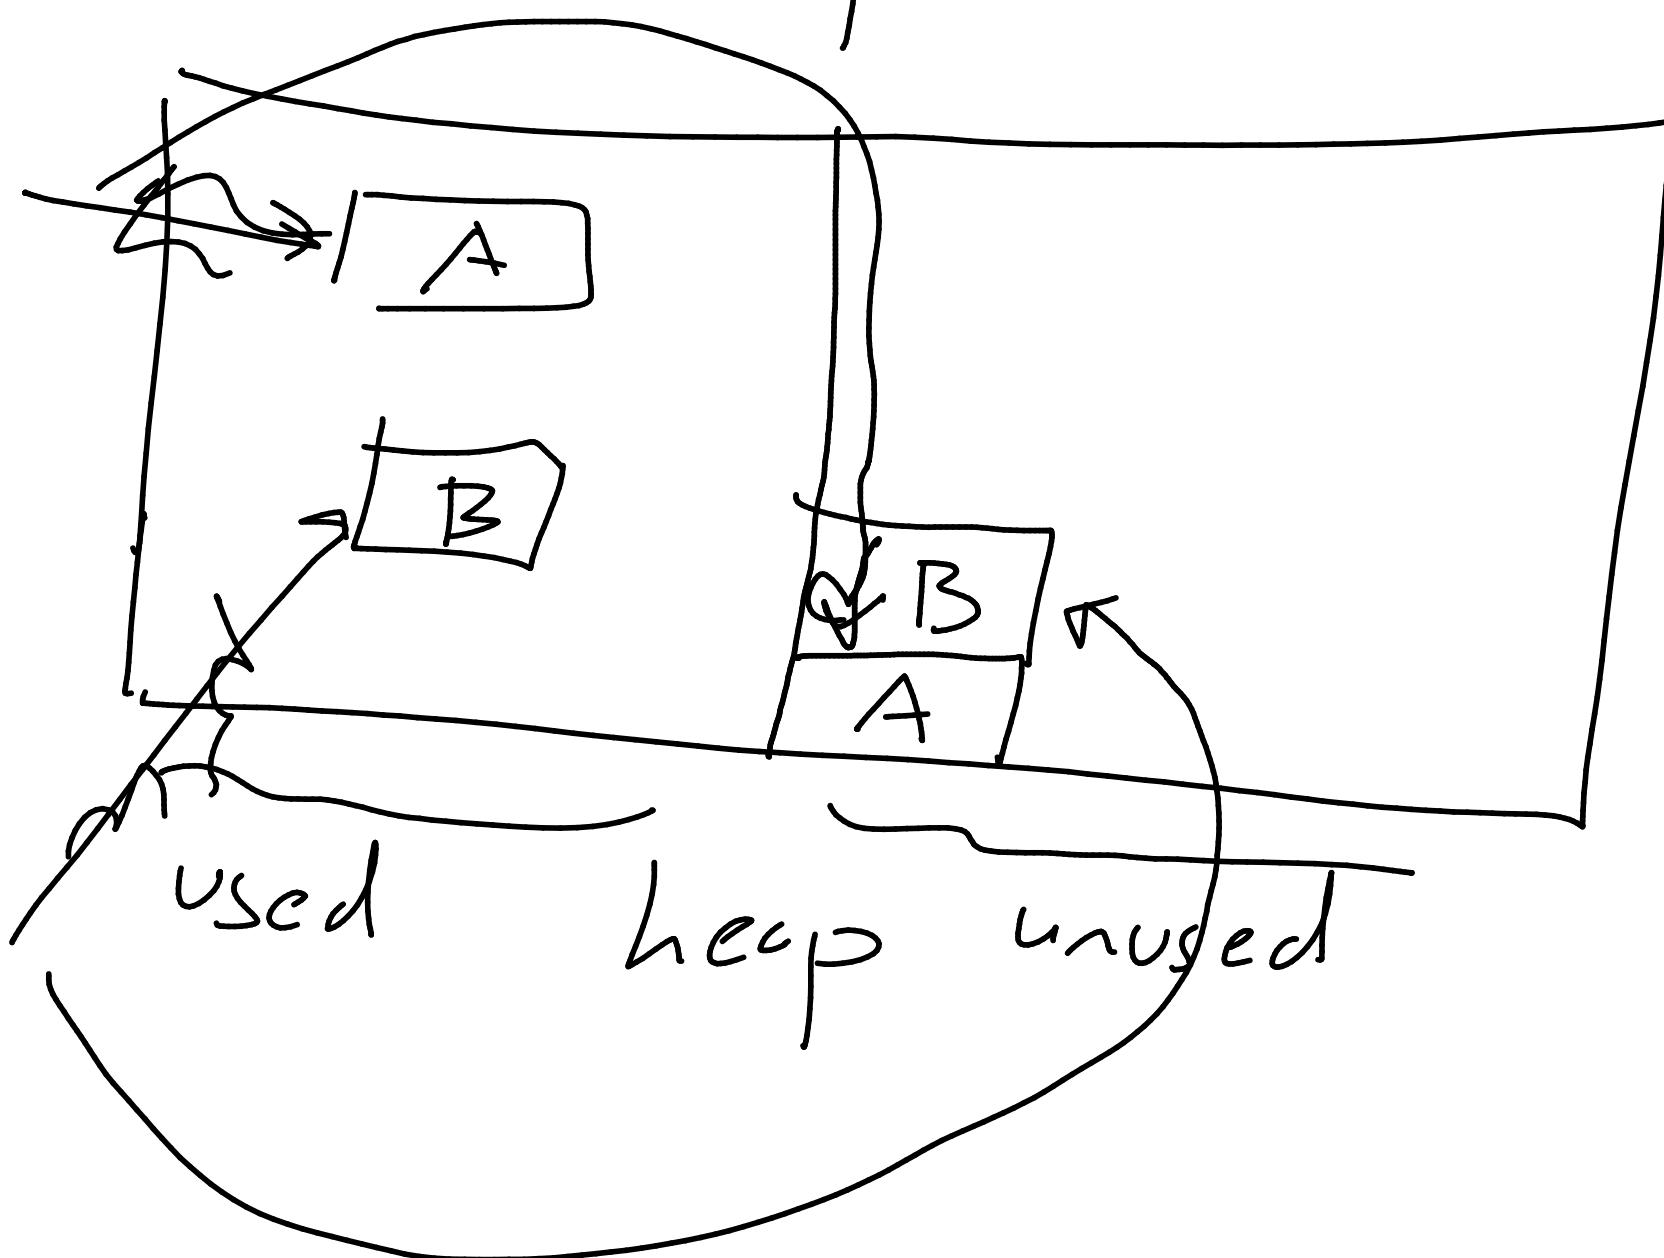
free

Our free memory is now broken
into smaller chunks (fragmented)
- limits the size of future objects
- and makes it slow to find
memory big enough for your
object

[would be just as bad if we
were freeing manually]

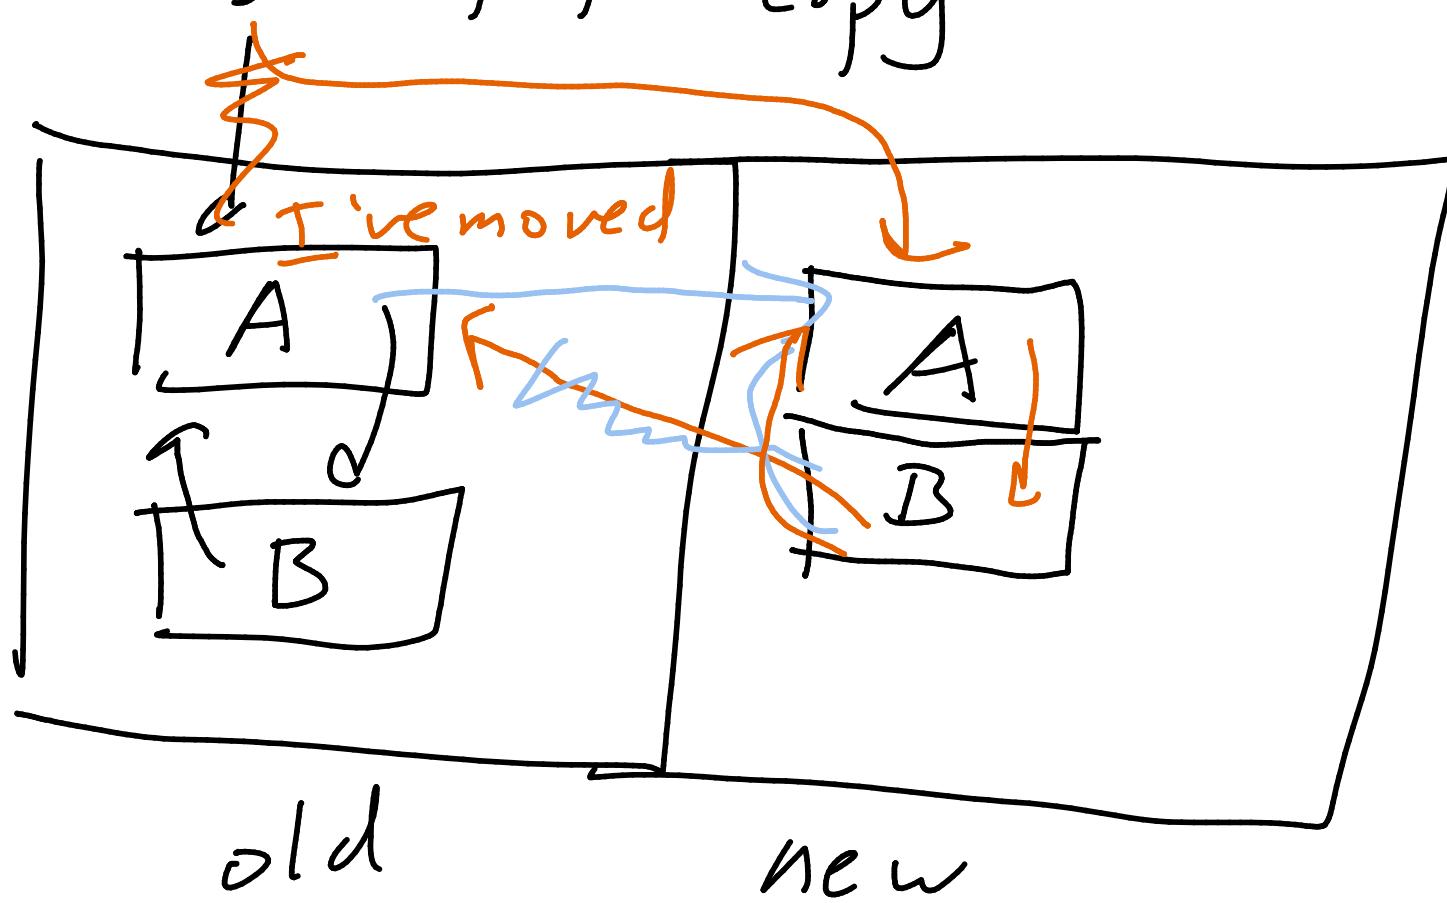
③ Copying (stop-and-copy)

- similar to mark and sweep
- difference: when you get to a live object, you copy it to an alternate heap



- copy each object to the unused heap, reset ptrs, and when done, the old heap is now the unused one.

whenever you copy an object from the old heap to the new, you write in the old location a "forwarding address" so that any other objects pointing there can update their ptrs if they try to copy



This works great, it results in defragmented memory, and uses twice as much memory to pull off.

When done, old heap is just completely done until you reverse roles next time

Takes a little more than mark/sweep
because it has to copy. (?)

But it doesn't have to free individual
chunks like mark/sweep does, so
takes less?

One hybrid approach that is
common (Java VM does this)

is generational gc

Studies have shown that young objects
are most likely to die

→ Use one gc technique more
frequently only on young objects